

# Assignment 2: Word2vec - Code Report

Student: Omar Adel Aly

Email: [omar.aly.ms@alexu.edu.eg](mailto:omar.aly.ms@alexu.edu.eg)

c)

```
class PartialParse(object):
    def __init__(self, sentence):
        """Initializes this partial parse.

        @param sentence (list of str): The sentence to be parsed as a list of words.
            Your code should not modify the sentence.

        """
        # The sentence being parsed is kept for bookkeeping purposes. Do NOT alter it
        # in your code.
        self.sentence = sentence

        ### YOUR CODE HERE (3 Lines)
        ### Your code should initialize the following fields:
        ###     self.stack: The current stack represented as a list with the top of the
        stack as the
        ###         Last element of the list.
        ###     self.buffer: The current buffer represented as a list with the first
        item on the
        ###         buffer as the first item of the list
        ###     self.dependencies: The list of dependencies produced so far.
        Represented as a list of
        ###         tuples where each tuple is of the form (head, dependent).
        ###         Order for this list doesn't matter.
        ###
        ### Note: The root token should be represented with the string "ROOT"
        ### Note: If you need to use the sentence object to initialize anything, make
        sure to not directly
        ###         reference the sentence object. That is, remember to NOT modify the
        sentence object.

        self.stack = ["ROOT"]
        self.buffer = sentence[:]
        self.dependencies = []

        ### END YOUR CODE
```

```

def parse_step(self, transition):
    """Performs a single parse step by applying the given transition to this
    partial parse

    @param transition (str): A string that equals "S", "LA", or "RA" representing
    the shift,
                                left-arc, and right-arc transitions. You can assume the
    provided
                                transition is a legal transition.

    """
    ### YOUR CODE HERE (~7-12 Lines)
    ### TODO:
    ###     Implement a single parsing step, i.e. the logic for the following as
    ###     described in the pdf handout:
    ###         1. Shift
    ###         2. Left Arc
    ###         3. Right Arc

    if transition == "S":
        self.stack.append(self.buffer.pop(0))
    elif transition == "LA":
        second_most_recent = self.stack.pop(-2)
        self.dependencies.append((self.stack[-1], second_most_recent))
    else:
        first_most_recent = self.stack.pop(-1)
        self.dependencies.append((self.stack[-1], first_most_recent))

    ### END YOUR CODE

```

d)

```

def minibatch_parse(sentences, model, batch_size):
    """Parses a list of sentences in minibatches using a model.

    @param sentences (list of list of str): A list of sentences to be parsed
                                            (each sentence is a list of words and each
    word is of type string)

    @param model (ParserModel): The model that makes parsing decisions. It is assumed
    to have a function
                                model.predict(partial_parses) that takes in a list of
    PartialParses as input and
                                returns a list of transitions predicted for each parse.
    That is, after calling
                                transitions = model.predict(partial_parses)
    transitions[i] will be the next transition to apply to
    partial_parses[i].

```

```

    @param batch_size (int): The number of PartialParses to include in each minibatch

    @return dependencies (list of dependency lists): A list where each element is the
dependencies
                                                    list for a parsed sentence.
Ordering should be the
                                                    same as in sentences (i.e.,
dependencies[i] should
                                                    contain the parse for
sentences[i]).
    """
    dependencies = []

    ### YOUR CODE HERE (~8-10 Lines)
    ### TODO:
    ###     Implement the minibatch parse algorithm. Note that the pseudocode for this
algorithm is given in the pdf handout.
    ###
    ###     Note: A shallow copy (as denoted in the PDF) can be made with the "=" sign
in python, e.g.
    ###
    ###         unfinished_parses = partial_parses[:].
    ###
    ###         Here `unfinished_parses` is a shallow copy of `partial_parses`.
    ###
    ###         In Python, a shallow copied list like `unfinished_parses` does not
contain new instances
    ###
    ###         of the object stored in `partial_parses`. Rather both lists refer
to the same objects.
    ###
    ###         In our case, `partial_parses` contains a list of partial parses.
`unfinished_parses`
    ###
    ###         contains references to the same objects. Thus, you should NOT use
the `del` operator
    ###
    ###         to remove objects from the `unfinished_parses` list. This will free
the underlying memory that
    ###
    ###         is being accessed by `partial_parses` and may cause your code to
crash.

    partial_parses = [PartialParse(sentence) for sentence in sentences]
    unfinished_parses = partial_parses

    while len(unfinished_parses) > 0:
        minibatch = unfinished_parses[0:batch_size]
        while len(minibatch) > 0:
            transitions = model.predict(minibatch)
            for i, parse in enumerate(minibatch):
                parse.parse_step(transitions[i])

```

```

        minibatch = [parse for parse in minibatch if len(parse.stack) > 1 or
len(parse.buffer) > 0]

        unfinished_parsers = unfinished_parsers[batch_size:]

        dependencies = []
        for n in range(len(sentences)):
            dependencies.append(partial_parsers[n].dependencies)
        ### END YOUR CODE

    return dependencies

```

e)

1)

```

#!/usr/bin/env python3
# -*- coding: utf-8 -*-
"""
CS224N 2023-2024: Homework 2
parser_model.py: Feed-Forward Neural Network for Dependency Parsing
Sahil Chopra <schopra8@stanford.edu>
Haoshen Hong <haoshen@stanford.edu>
"""

import argparse
import numpy as np

import torch
import torch.nn as nn
import torch.nn.functional as F

class ParserModel(nn.Module):
    """ Feedforward neural network with an embedding layer and two hidden layers.
    The ParserModel will predict which transition should be applied to a
    given partial parse configuration.

    PyTorch Notes:
        - Note that "ParserModel" is a subclass of the "nn.Module" class. In PyTorch
all neural networks
        are a subclass of this "nn.Module".
        - The "__init__" method is where you define all the layers and parameters
(embedding layers, linear layers, dropout layers, etc.).
        - "__init__" gets automatically called when you create a new instance of your
class, e.g.
            when you write "m = ParserModel()".

```

- Other methods of `ParserModel` can access variables that have `"self."` prefix.

Thus,

you should add the `"self."` prefix layers, values, etc. that you want to

utilize

in other `ParserModel` methods.

- For further documentation on `"nn.Module"` please see

<https://pytorch.org/docs/stable/nn.html>.

"""

```
def __init__(self, embeddings, n_features=36,
             hidden_size=200, n_classes=3, dropout_prob=0.5):
```

""" Initialize the parser model.

@param embeddings (ndarray): word embeddings (num\_words, embedding\_size)

@param n\_features (int): number of input features

@param hidden\_size (int): number of hidden units

@param n\_classes (int): number of output classes

@param dropout\_prob (float): dropout probability

"""

```
super(ParserModel, self).__init__()
```

```
self.n_features = n_features
```

```
self.n_classes = n_classes
```

```
self.dropout_prob = dropout_prob
```

```
self.embed_size = embeddings.shape[1]
```

```
self.hidden_size = hidden_size
```

```
self.embeddings = nn.Parameter(torch.tensor(embeddings))
```

```
### YOUR CODE HERE (~9-10 Lines)
```

```
### TODO:
```

```
### 1) Declare `self.embed_to_hidden_weight` and
```

```
`self.embed_to_hidden_bias` as `nn.Parameter`.
```

```
### Initialize weight with the `nn.init.xavier_uniform` function and
```

```
bias with `nn.init.uniform`
```

```
### with default parameters.
```

```
### 2) Construct `self.dropout` layer.
```

```
### 3) Declare `self.hidden_to_logits_weight` and
```

```
`self.hidden_to_logits_bias` as `nn.Parameter`.
```

```
### Initialize weight with the `nn.init.xavier_uniform` function and
```

```
bias with `nn.init.uniform`
```

```
### with default parameters.
```

```
###
```

```
### Note: Trainable variables are declared as `nn.Parameter` which is a  
commonly used API
```

```
### to include a tensor into a computational graph to support updating  
w.r.t its gradient.
```

```

    """ Here, we use Xavier Uniform Initialization for our Weight
initialization.

    """ It has been shown empirically, that this provides better initial
weights

    """ for training networks than random uniform initialization.
    """ For more details checkout this great blogpost:
    """
http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization
    """

    """ Please see the following docs for support:
    """ nn.Parameter: https://pytorch.org/docs/stable/nn.html#parameters
    """ Initialization: https://pytorch.org/docs/stable/nn.init.html
    """ Dropout: https://pytorch.org/docs/stable/nn.html#dropout-layers
    """
    """ See the PDF for hints.

    self.embed_to_hidden_weight = nn.Parameter(torch.empty(n_features *
self.embed_size, hidden_size))
    self.embed_to_hidden_bias = nn.Parameter(torch.empty(hidden_size))
    nn.init.xavier_uniform_(self.embed_to_hidden_weight)
    nn.init.uniform_(self.embed_to_hidden_bias)

    self.dropout = nn.Dropout(p=dropout_prob)

    self.hidden_to_logits_weight = nn.Parameter(torch.empty(hidden_size,
n_classes))
    self.hidden_to_logits_bias = nn.Parameter(torch.empty(n_classes))
    nn.init.xavier_uniform_(self.hidden_to_logits_weight)
    nn.init.uniform_(self.hidden_to_logits_bias)

    """ END YOUR CODE

def embedding_lookup(self, w):
    """ Utilize `w` to select embeddings from embedding matrix `self.embeddings`
    @param w (Tensor): input tensor of word indices (batch_size, n_features)

    @return x (Tensor): tensor of embeddings for words represented in w
    (batch_size, n_features * embed_size)
    """

    """ YOUR CODE HERE (~1-4 Lines)
    """ TODO:
    """ 1) For each index `i` in `w`, select `i`th vector from self.embeddings
    """ 2) Reshape the tensor using `view` function if necessary
    """

```

```
    ### Note: All embedding vectors are stacked and stored as a matrix. The model receives  
    ### a list of indices representing a sequence of words, then it calls this lookup  
    ### function to map indices to sequence of embeddings.  
    ###  
    ### This problem aims to test your understanding of embedding lookup,  
    ### so DO NOT use any high level API like nn.Embedding  
    ### (we are asking you to implement that!). Pay attention to tensor shapes  
    ### and reshape if necessary. Make sure you know each tensor's shape before you run the code!
```

```
    ###  
    ### Pytorch has some useful APIs for you, and you can use either one  
    ### in this problem (except nn.Embedding). These docs might be helpful:  
    ### Index select:  
https://pytorch.org/docs/stable/torch.html#torch.index\_select  
    ### Gather: https://pytorch.org/docs/stable/torch.html#torch.gather  
    ### View: https://pytorch.org/docs/stable/tensors.html#torch.Tensor.view  
    ### Flatten: https://pytorch.org/docs/stable/generated/torch.flatten.html
```

```
    batch_size, n_features = w.shape  
    embed_size = self.embeddings.shape[1]  
    x = torch.index_select(self.embeddings, 0, w.flatten())  
    x = x.view(batch_size, n_features * embed_size)
```

```
    ### END YOUR CODE  
    return x
```

```
def forward(self, w):  
    """ Run the model forward.
```

*Note that we will not apply the softmax function here because it is included in the loss function nn.CrossEntropyLoss*

*PyTorch Notes:*

- Every nn.Module object (PyTorch model) has a `forward` function.*
- When you apply your nn.Module to an input tensor `w` this function is applied to the tensor.*

*For example, if you created an instance of your ParserModel and applied it to some `w` as follows,*

*the `forward` function would be called on `w` and the result would be stored in the `output` variable:*

```
    model = ParserModel()
```

```

        output = model(w) # this calls the forward function
        - For more details checkout:
https://pytorch.org/docs/stable/nn.html#torch.nn.Module.forward

        @param w (Tensor): input tensor of tokens (batch_size, n_features)

        @return logits (Tensor): tensor of predictions (output after applying the
Layers of the network)

                                without applying softmax (batch_size, n_classes)
    """
    ### YOUR CODE HERE (~3-5 lines)
    ### TODO:
    ###     Complete the forward computation as described in write-up. In addition,
include a dropout layer
    ###     as declared in `__init__` after ReLU function.
    ###
    ### Note: We do not apply the softmax to the logits here, because
    ### the loss function (torch.nn.CrossEntropyLoss) applies it more efficiently.
    ###
    ### Please see the following docs for support:
    ###     Matrix product: https://pytorch.org/docs/stable/torch.html#torch.matmul
    ###     ReLU:
https://pytorch.org/docs/stable/nn.html?highlight=relu#torch.nn.functional.relu

    x = self.embedding_lookup(w)
    h = F.relu(torch.matmul(x, self.embed_to_hidden_weight) +
self.embed_to_hidden_bias)
    h_drop = self.dropout(h)
    logits = torch.matmul(h_drop, self.hidden_to_logits_weight) +
self.hidden_to_logits_bias

    ### END YOUR CODE
    return logits

```

2)

```

# -----
# Primary Functions
# -----
def train(parser, train_data, dev_data, output_path, batch_size=1024, n_epochs=10,
lr=0.0005):
    """ Train the neural dependency parser.

    @param parser (Parser): Neural Dependency Parser
    @param train_data ():
    @param dev_data ():

```



```

@param output_path (str): Path to which model weights and results are written.
@param batch_size (int): Number of examples in a single batch
@param n_epochs (int): Number of training epochs
@param lr (float): Learning rate
"""

best_dev_UAS = 0

### YOUR CODE HERE (~2-7 lines)
### TODO:
###     1) Construct Adam Optimizer in variable `optimizer`
###     2) Construct the Cross Entropy Loss Function in variable `loss_func` with
`mean`
###         reduction (default)
###
### Hint: Use `parser.model.parameters()` to pass optimizer
###         necessary parameters to tune.
### Please see the following docs for support:
###     Adam Optimizer: https://pytorch.org/docs/stable/optim.html
###     Cross Entropy Loss:
https://pytorch.org/docs/stable/nn.html#crossentropyloss

optimizer = optim.Adam(parser.model.parameters(), lr=lr)
loss_func = nn.CrossEntropyLoss()

### END YOUR CODE

for epoch in range(n_epochs):
    print("Epoch {:} out of {:}".format(epoch + 1, n_epochs))
    dev_UAS = train_for_epoch(parser, train_data, dev_data, optimizer, loss_func,
batch_size)
    if dev_UAS > best_dev_UAS:
        best_dev_UAS = dev_UAS
        print("New best dev UAS! Saving model.")
        torch.save(parser.model.state_dict(), output_path)
    print("")

def train_for_epoch(parser, train_data, dev_data, optimizer, loss_func, batch_size):
    """ Train the neural dependency parser for single epoch.

    Note: In PyTorch we can signify train versus test and automatically have
    the Dropout Layer applied and removed, accordingly, by specifying
    whether we are training, `model.train()`, or evaluating, `model.eval()`

```

```

@param parser (Parser): Neural Dependency Parser
@param train_data ():
@param dev_data ():
@param optimizer (nn.Optimizer): Adam Optimizer
@param loss_func (nn.CrossEntropyLoss): Cross Entropy Loss Function
@param batch_size (int): batch size

@return dev_UAS (float): Unlabeled Attachment Score (UAS) for dev data
"""

parser.model.train() # Places model in "train" mode, i.e. apply dropout layer
n_minibatches = math.ceil(len(train_data) / batch_size)
loss_meter = AverageMeter()

with tqdm(total=n_minibatches) as prog:
    for i, (train_x, train_y) in enumerate(minibatches(train_data, batch_size)):
        optimizer.zero_grad() # remove any baggage in the optimizer
        loss = 0. # store loss for this batch here
        train_x = torch.from_numpy(train_x).long()
        train_y = torch.from_numpy(train_y.nonzero()[1]).long()

        ### YOUR CODE HERE (~4-10 lines)
        ### TODO:
        ###      1) Run train_x forward through model to produce `logits`
        ###      2) Use the `loss_func` parameter to apply the PyTorch
CrossEntropyLoss function.
        ###      This will take `logits` and `train_y` as inputs. It will output
the CrossEntropyLoss
        ###      between softmax(`logits`) and `train_y`. Remember that
softmax(`logits`)
        ###      are the predictions ( $y^{\wedge}$  from the PDF).
        ###      3) Backprop losses
        ###      4) Take step with the optimizer
        ### Please see the following docs for support:
        ###      Optimizer Step:
https://pytorch.org/docs/stable/optim.html#optimizer-step

        logits = parser.model(train_x)
        loss = loss_func(logits, train_y)
        loss.backward()
        optimizer.step()

        ### END YOUR CODE
        prog.update(1)
        loss_meter.update(loss.item())

```

```

print ("Average Train Loss: {}".format(loss_meter.avg))

print("Evaluating on dev set",)
parser.model.eval() # Places model in "eval" mode, i.e. don't apply dropout layer
dev_UAS, _ = parser.parse(dev_data)
print("- dev UAS: {:.2f}".format(dev_UAS * 100.0))
return dev_UAS

```

# -----

## Result

```
=====
TESTING
=====
```

```

Restoring the best model weights found on the dev set
D:\2025\github\nlp-assignments\2-word2vec\student\run.py:160:
ring unpickling (See https://github.com/pytorch/pytorch/blob/main/
onger be allowed to be loaded via this mode unless they are ex
r any issues related to this experimental feature.
  parser.model.load_state_dict(torch.load(output_path))
Final evaluation on test set
2919736it [00:00, 116782312.17it/s]
- test UAS: 89.28
Done!

```

- test UAS: 89.28